

Design Document for Multi-user Infinite Canvas Thingy (MICT)

Don Huckle, Ben Kaplan, Mark Wyrzykowski, Rob Wiesler

October 4, 2010

Contents

1	Version Info	2
2	Tools	2
2.1	Serializing the Tools	2
2.2	Mouse Events	2
2.3	Drawing to a Graphics Context	2
2.4	Tool Properties	3
3	Client Program Specification	3
3.1	Client	3
3.2	ClientState	3
3.3	Toolbox	4
3.3.1	ToolButtons	4
3.4	Canvas Viewport	4
3.5	Java-Python Bridge	4
3.6	Canvas Control Panel	4
3.7	Launching the MICT Client	6
4	Client-Server Specification	6
4.1	Connecting to the Server	6
4.2	Client-Side Generated Communication	7
4.3	Server Response to User Edits	8
4.4	Server storage of Canvas Chunks	8
4.5	Server-Side Generated Communication	9
5	Appendix: UML Dependency Diagram	10
6	Bibliography	10

1 Version Info

Version	Info	Date
v0.1	Initial Draft	10/4/2010

2 Tools

The Tools are the heavy lifters of the program. Each Tool is self-contained and is responsible for drawing itself on at least three kinds of graphics contexts:

- The context of the user's viewport
- The context of the infinite canvas, as it exists on the server
- The context of the viewports of other users who are viewing a particular section of canvas

2.1 Serializing the Tools

Each of the Tools will be stored on the server and transmitted to the Client upon log-in. To facilitate this, each of the Tools should be implemented in Jython. Preliminary research shows that the Jython classes can be serialized using Java's ObjectOutputStream, but this has not yet confirmed satisfactorily.

2.2 Mouse Events

The Tool is expected to implement MousePressed, MouseDragged, and MouseReleased methods that are called when the mouse is pressed, moved (while pressed), and then released while this Tool is active. Each of these methods returns a String, which will be transmitted to the server in order to apply the change to the global canvas. These methods are passed the canvas's Graphics object and the Point on the canvas where the event was fired from.

2.3 Drawing to a Graphics Context

Here, the Tool will be passed a String generated by the tool and a Graphics object to draw the figure on. The changes to the canvas generated by this method should be identical to the changes generated by the tool on the client side. To this end, Tools acting on a user's local viewport will only draw to the local graphics context based on the same String passed to the server-side Tool incarnation.

2.4 Tool Properties

Each tool will need to know its name in both human-readable and shortened internal forms, and an icon. The icon and human-readable name (presented as a tooltip) will be used on the client-side to represent the tool to the user.

3 Client Program Specification

This section describes the layout of the client program. It will discuss all the components of that program and how they interact with each other. There are five main parts to client-side architecture:

- the Client
- the ClientState
- the Toolbox
- the Canvas Control Panel
- the Canvas Viewport

3.1 Client

The root of the client program is the Client class. This is the application/applet itself. Other than initializing the other components and coordinating communication between them, it is also responsible for setting the location of the canvas, either as a result of initially connecting to the server, through a jump command, or by panning the viewport. The Client will hold a ClientState object that stores the shared settings. In addition, the GUI will consist of three components:

- the Toolbox
- the Canvas Control Panel
- the Canvas Viewport

3.2 ClientState

A single ClientState object will be initialized for each instance of the program. The ClientState object will be used to store the shared state. Any field that needs to be accessed and modified by multiple classes should be placed in here, including but not limited to the currently selected tool and color, the current location of the canvas, and the persistent connection to the server. The use of a singleton structure for shared state will reduce the coupling between all the components, slightly increasing complexity, but opening up the client-side architecture for potential future improvements, such as multiple viewports and simultaneous connections to multiple servers.

3.3 Toolbox

The Toolbox will contain a series of Tools. Each Tool (which extends the abstract class `mict.tools.Tool`) shall be defined in Jython so that we can serialize the class and transmit it from the server to the client. This way, every client will support every tool the server has without any prior configuration on the part of the user. Additionally, this provides us with an opportunity to leverage a more consistent Tool architecture than might exist should both a client- and server-side implementation exist for each Tool.

3.3.1 ToolButtons

For each Tool, the Client will make a ToolButton. ToolButton is a subclass of JButton. It has an ActionListener that sets the associated tool as the active tool within ClientState when the button is pressed. Each ToolButton will set its icon and tooltip as directed with the underlying Tool class.

3.4 Canvas Viewport

The Canvas Viewport is a JPanel embedded in the Client UI that will display a section of the overall canvas. The user will draw on the Viewport with a Tool in order to modify the canvas, and other users modifications will be updated here.

The Canvas shall have a single MouseListener, which will wait for MousePressed, MouseDragged, and MouseReleased events and dispatch those events to the currently selected tool. As these methods are called, the Canvas shall transmit the serialized commands to the server via the Client Connection, which will be discussed later in this document.

3.5 Java-Python Bridge

There will be a single Java class and a single Python module (`JythonBridge.java` and `javabridge.py`) to facilitate communication between Jython and Java components. Any Java component that needs to access something from Jython should go through JythonBridge. Because of the way Jython works, Python components can access the Java libraries directly if they need to extend them (for instance, extending Tools). However, if they require an instance of a Java object, they should go through javabridge.

3.6 Canvas Control Panel

Apart from normal editing actions, users with elevated privileges will have access to a separate tab in the Toolbox. In this tab there will be additional functions available based on the level of the user. Among these functions are

the ability to lock parts of the canvas, kick users, move other users around the canvas, ban/unban users, and modify the permissions of other users.

When a user locks an area, they gain power over that area and are the only one (Administrators not included) that can edit it unless they decide to invite other users to edit the area. The administration panel is context-aware, so options present in the Control Panel deactivate when the user leaves the area in which he or she has elevated privileges, and re-activate upon return.

Administrators have this power over the entire canvas. They can lock and unlock sections even if that section had been locked by another user.

In addition to modifying privileges for the canvas, Administrators can reconfigure the canvas in real time, and can save the canvas and stop the server gracefully.

- The options in this panel will be implemented as Tools if possible, and will only be dealt with as a special case if the Tool interface does not provide all of the required functionality of the option.
- These controls change as the user moves over the canvas. The server will communicate to the Client which sections are valid options for the user given the user's current location on the canvas. For instance, the server may be configured such that a particular user owns reserved areas of the canvas. Naturally in this scenario, the administrative controls for reserved segments will be contextuallized by the position of the user on the canvas. Therefore, in a public area, the administrative controls of a user of a locked area somewhere else will be disabled in the Canvas Control Panel, and the network protocol will stop responding to administrative actions.
- Users will only be presented with options appropriate to their permissions. If they do not have permission to acquire sections for instance, they will not see that button presented to them. If they have permission but are unable to acquire it (for instance, if someone else owns it), the button will be inactive.
- Administrator-level users will have available to them in the Control Panel options to trigger the Canvas.User.Kick, Canvas.User.Ban, and Canvas.User.Pardon functions from section 3.4 of the requirements document.

3.7 Launching the MICT Client

When the program is launched as an application, the Toolbox will be empty and the Canvas Viewport will be replaced by two JTextFields, a JPasswordField, and a JButton. The first JTextField will be used to enter the server to connect to. The second will be the username, and the JPasswordField will be used to enter the password. The JButton will tell the client to connect to the server specified with the given log-in credentials. If login is successful, the fields will be replaced with the Canvas Viewport. This implements the Canvas.Connect requirement mentioned in section 3.1 of the SRS document.

In the case that the client is launched as an applet, the above behavior will occur, with the exception that the presence of an HTML applet parameter providing the name of the server will result in the absence of a text area used to specify the server.

- When the user presses the log-in button, the ClientConnection will be created. Upon successful connection, the Client will create the ClientState object, retrieve the Tools from ther server and create the ToolButtons, and then add the Canvas Viewport and set the associated Graphics in the ClientState.
- Depending on the user's permissions and location, the Canvas Control panel may also be populated to hold the interface for administrative actions specified in Section 3.3 of the Software Requirements Document.

4 Client-Server Specification

The client server interaction is based on a TCP connection between the server and the client. In per-user communication there are three major threads driving interaction:

- the server's Waiter thread
- the client's ClientConnection thread
- the client's GUI thread

4.1 Connecting to the Server

Upon clicking the "connect" button at the start of the program execution, the client GUI thread will generate a ClientConnection object responsible for maintaining a persistent connection with the server. On the server side exists a corresponding Waiter object that represents the user within the architecture of the server. After the initial SSL handshaking, the client will

transmit its username and a secure hash of its password and the hostname. This is for security reasons as follows:

- The password must be hashed at least on disk to limit the risk of revealing the user's password to persons with hard or soft access to the server.
- The password must also be hashed before being handed over to the server to avoid giving the server administrator a password potentially in use on other servers and services.
- Since a malicious administrator with access to the hashed version of a password may potentially modify the client to send a hashed password to another server, thus gaining admission, the password must be hashed along with server-specific information in order for the hashed password to maintain validity for only a single server.
- The Internet-facing hostname of the server¹ is ideally suited for this task. For one, if a server-wide arbitrary string were given to the user as the server-specific string, it would be simple to momentarily change the server-specific string on a malicious server to match the one used on another server, rendering the measure useless. However, the relatively constant hostname is not so easily spoofed by a malicious administrator, since the client is aware of its true value even before connecting².

If the authentication provided by the client is accepted by the server, the server will respond with a Message of the Day (MotD), and an initial set of coordinates. The client should respond with an optional request to change location, and notify the server of the desired width and height of its viewport, which the server should respond to with an initial render of the canvas at that point, sized to fit the viewport.

4.2 Client-Side Generated Communication

When an event is triggered on the client side, it will often result in a message being sent to the `ClientConnection` object. This object will then package the information in a protocol wrapper, and pass it on to the server. To prevent network communications from locking up server- or client-side logic, both the `ClientConnection` and `Waiter` objects should handle incoming messages in a separate thread. Operations that will cause communication to be generated by the client include `mouse_clicked`, `mouse_dragged`, and `mouse_released` events. These events have event handlers associated with them that instruct the active Tool to encode the user's current actions into a String, which the client then packages with the internal name of the Tool and sends to the `Waiter` object on the server side.

¹As in `rdebase.case.edu` as opposed to `rdebase`

²Except notably in the case of Applet-based deployments

4.3 Server Response to User Edits

When the Waiter object gets an instruction to apply a tool to the canvas, it pushes the request, along with the user's name and viewing rectangle, down-architecture to the server's CanvasManager object. The purpose of the Canvas Manager is to handle the intricacies of the infinite canvas abstraction. The Canvas Manager retrieves the tool corresponding to the tool name in the message from the ToolManager entity, and locates the Chunk object closest to the center of the user's viewport. It then creates a ChunkGraphics graphics context for that chunk, which performs three functions:

- It handles the translation from the user's viewport coordinate system into the coordinate system of the infinite canvas, using offsets provided by the Waiter object's bounding rectangle field.
- It keeps track of drawing operations that extend over the edge of the Chunk's boundaries, and applies those draw operations³ to the Chunk objects that exist beyond the borders of the initial Chunk object.
- It prevents Tools wielded by a user without permission to write in an area from modifying said area, which may or may not have been delimited by Chunk boundaries. For instance, if user Bob controls a smallish circular area and Alice does not have write permissions within that area, a line drawn by Alice through Bob's circular area and extending some distance to either side will resolve on the canvas on the outside of Bob's area, but on inside. In response to an illegal write attempt, it passes a message back to the client, which notifies the user in a subtle way that something went wrong, and eventually updates the user with a correct version of the visible canvas.

If there are other users active in the region of Chunks that were modified by the Tool, the Canvas Manager locates them among the list of Waiter objects and pushes updates to the affected areas out to them, with care to only update the intersection of the affected areas and the views of the other users.

4.4 Server storage of Canvas Chunks

Since an infinite⁴ canvas will not necessarily fit into memory all at once, a backend layer that consists of the CanvasManager and the DatabaseLayer objects performs some abstraction of the canvas into Chunks, as mentioned before, and stores only a limited amount of them in memory at once. The maximum number of Chunks in memory is governed by the MAXCHUNKS field of the CanvasManager class. When the Canvas Manager gets a request for a

³with appropriate offsets, using another ChunkGraphics context

⁴or, at least arbitrarily large

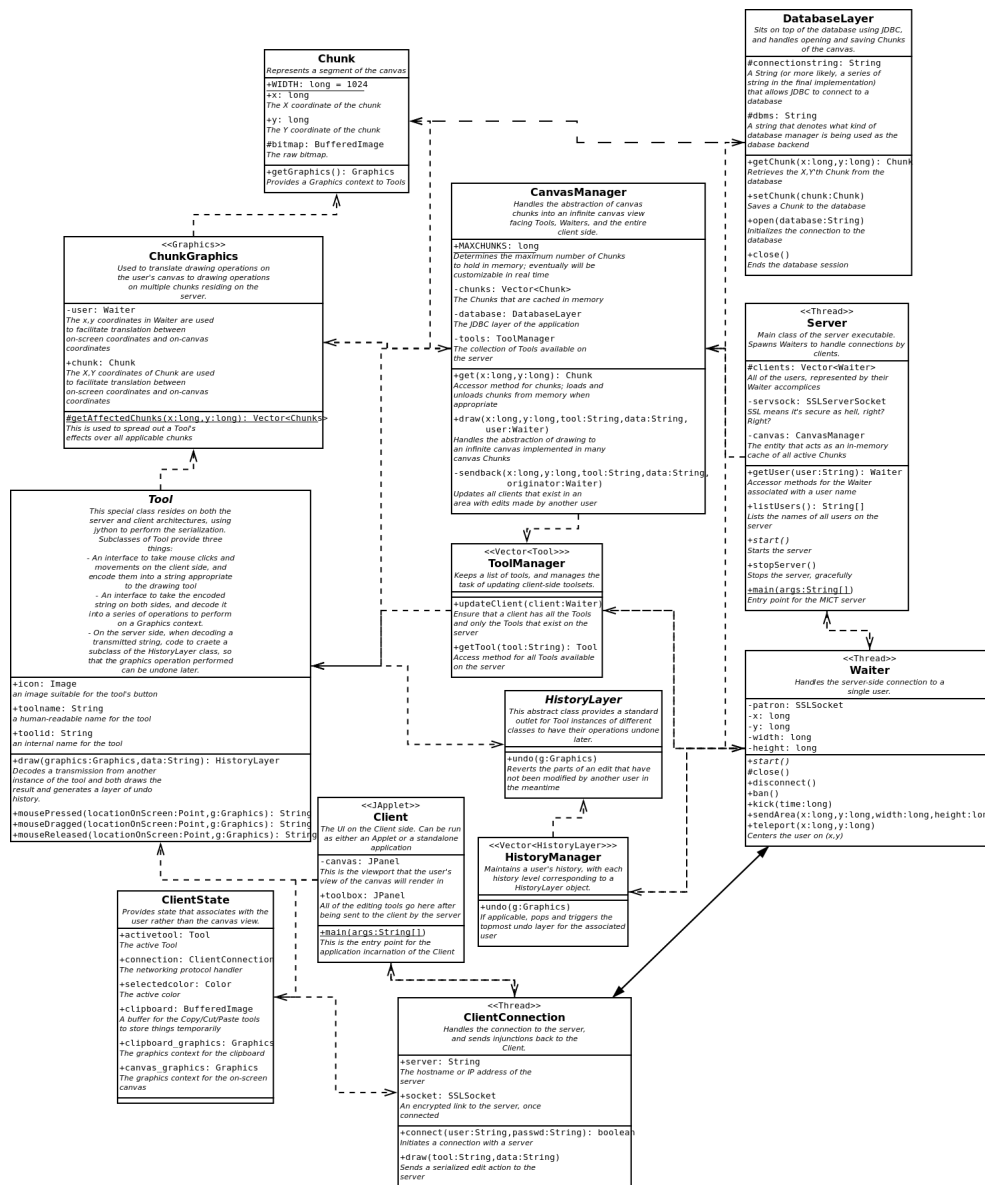
Chunk, it servers its cache for the matching Chunk. If it finds it, it returns it. Otherwise, it requests the Chunk from the Database Layer. If it needs to, it passes inactive Chunks back onto disk by returning them to the Database Layer, which saves them to the database.

In case of catastrophic system failure, a thread internal to the Canvas Manager periodically passes all Chunks back to the Database Layer to save them.

4.5 Server-Side Generated Communication

When the Server has some action that it needs the client to take, such as updating the client's canvas with edits made by another user, the server will send an injunction to the client specifying either an update to the client's canvas or a message to hand to the user. The client connection Because Java's Swing architecture is not thread safe, the client will invoke the operation in a roundabout manner via `SwingUtilities.invokeLater(Runnable)`.

5 Appendix: UML Depemdency Diagram



6 Bibliography

Huckle, Donald, et al. Multi-User Infinite Canvas Thingy Vision and Scope Document. http://github.com/halcyone1024/MICT/raw/master/doc/vision_and_scope.pdf

Huckle, Donald, et al. Multi-User Infinite Canvas Thingy Software Requirement Specification.

http://github.com/halcyone1024/MICT/raw/master/doc/software_requirement_specification.pdf